

# **UNIX**

# **Inter-Process**

# **Communication**

Pipes, Messages,  
Semaphores

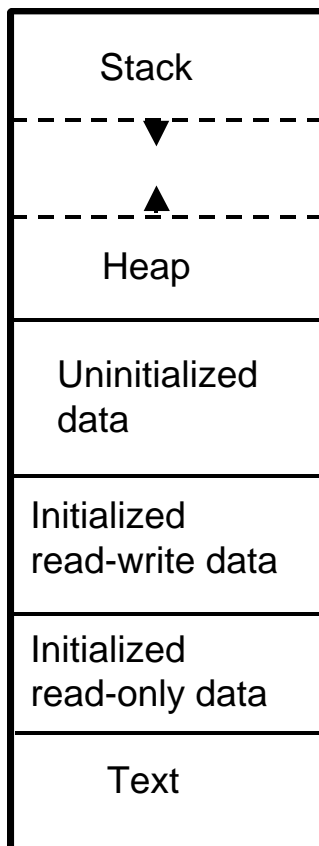
Prof. Badri R. Narayanan




# UNIX Process Model

- A *process* is an instance of a program in execution.
- Only way to create a new process is the *fork* system call.
- Each process has a *user context* and a *kernel context*.

User Context



Kernel Context



from file on *exec*.

# Relevant Attributes of Kernel Context of a Process

<i>Attribute</i>	<i>System Call</i>
Process Id	getpid()
Parent PID	getppid()
Real User Id	getuid()
Real Group Id	getgid()
Effective UID	geteuid() <i>set-user-ID</i>
Effective GID	getegid() <i>set-group-ID</i>

# UNIX Signals (Subset)

*Signals* notify processes of events.

- ➡ from a process
- ➡ from kernel

<i>Name</i>	<i>Description</i>	<i>Default Action</i>
SIGALRM	Alarm Clock	Terminate
SIGBUS	Bus Error	X with core
SIGCLD	Death of Child Process	Discarded
SIGCONT	Continue after SIGSTOP	Discarded
SIGHUP	Hangup	Terminate
SIGINT	Interrupt character	Terminate
SIGIO	I/O poss. on file desc.	Discarded
SIGKILL	Kill	Terminate
SIGPIPE	Write on pipe without reader	Terminate
SIGQUIT	Quit character	X with core
SIGSTOP	Stop	Stop (suspend)
SIGTSTP	Stop from keyboard	Stop
SIGURG	Urgent cond. on socket	Discarded
SIGUSR1	User defined signal	Terminate
SIGUSR2	User defined signal	Terminate

# Generating Signals

Signals are sent in response to conditions.

*Note:* Signals usually occur asynchronously.

<i>Condition</i>	<i>Example</i>
<i>kill</i> system call	int kill(int <i>pid</i> , int <i>sig</i> ) Send <i>sig</i> to <i>pid</i> . <i>Many variations.</i>
<i>kill</i> command	Issues <i>kill</i> system call
Terminal characters	Control-C -- SIGINT Control-Z -- SIGSTP
Hardware Conditions	Floating Point Error -- SIGFPE
Software Conditions	Out-of-band data on socket -- SIGURG

# Handling Signals

- Process can *catch* signals by providing signal handler functions.
- Process can choose to ignore all signals (except SIGKILL which guarantees termination).
- Process can allow the default action to occur.


## Using *signal* system call

```
#include <signal.h>
int (*signal (int sig, void (*func) (int))) (int);

/* SIG_IGN and SIG_DFL */
signal(SIGUSR1, SIG_IGN);
/* User-defined handler */
extern void myintr();
if (signal(SIGINT, SIG_IGN) != SIG_IGN)
    signal(SIGINT, myintr);
```

# Process Control

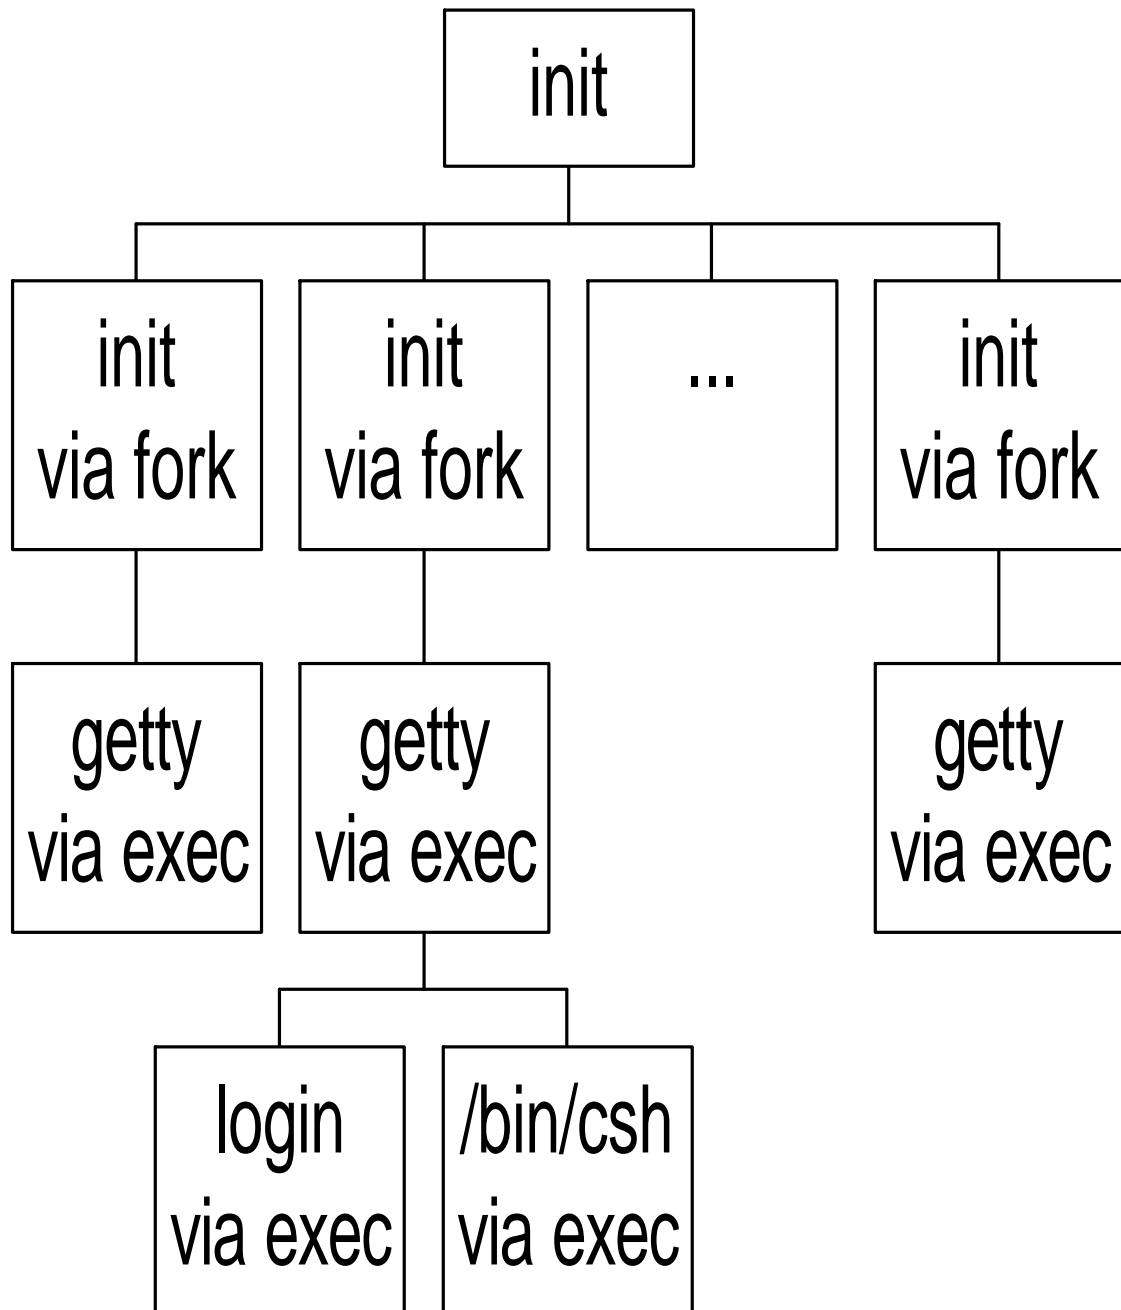
## *fork, exit, exec, wait*



- *fork* is called by parent process.
- It creates an identical child process sharing text, open file handles.
- *fork* returns child pid to parent.
- *fork* returns 0 to child process.

```
main()
{
    int childpid;
    if ( (childpid = fork()) == -1 ) {
        perror( "fork failed!");
        exit(1);
    } else
    if (childpid == 0) { /* child process */
    }
    else { /* parent process */
    }
}
```


# Process Relationships






# IPC Basics

## IPC on single system



## IPC across networked systems



# Need for Mutual Exclusion Revisited

Unique sequence number for print job

- Read shared sequence number file
- Use number from file
- Write incremented value to file

```
#define SEQFILE "jobnumber"
#define MAXBUFF 64
main() {
    int fdes, i, n, pid, seqno;
    char buffer[MAXBUFF + 1];

    pid = getpid();
    if ((fdes = open(SEQFILE, 2)) < 0)
        err_sys("Could not open %s", SEQFILE);
    for (i=0; i<5; i++) {
        my_lock(fdes);
        lseek(fdes, 0L, 0);
        if ((n = read(fdes, buff, MAXBUFF)) <= 0)
            err_sys("Reading error");
        buff[n] = '\0';
        if ((n = sscanf(buff, "%d\n", &seqno)) != 1)
            err_sys("sscanf error");
        printf("pid = %d, seq# = %d\n", pid, seqno);
        seqno++;
        sprintf(buff, "%03d\n", seqno);
        n = strlen(buff);
        lseek(fdes, 0L, 0);
        if (write(fdes, buff, n) != n)
            err_sys("Writing error");
        my_unlock(fdes);
    }
}
```

# File Locking for Mutual Exclusion

## Without Locks

```
my_lock(fdes)
int fdes;
{ return; }

my_unlock(fdes)
int fdes;
{ return; }
```

*Run program twice*  
a.out & a.out &

## Possible output

```
pid = 12, seq# = 1
pid = 12, seq# = 2

pid = 13, seq# = 2
pid = 13, seq# = 3
pid = 13, seq# = 4

pid = 12, seq# = 3

pid = 13, seq# = 4
```

File locking allows for safe and exclusive access of files.

## With Locks

```
#include <sys/file.h>
my_lock(fdes)
int fdes;
{ if (flock(fdes, LOCK_EX) == -1)
    err_sys("Cannot LOCK_EX");
}

my_unlock(fdes)
int fdes;
{ if (flock(fdes, LOCK_UN) == -1)
    err_sys("Cannot LOCK_UN");
}
```

## Possible output

```
pid = 15, seq# = 1
pid = 15, seq# = 2


pid = 16, seq# = 3
pid = 16, seq# = 4
pid = 16, seq# = 5

pid = 15, seq# = 6
pid = 15, seq# = 7

pid = 16, seq# = 8
pid = 16, seq# = 9
```

# Client-Server Example

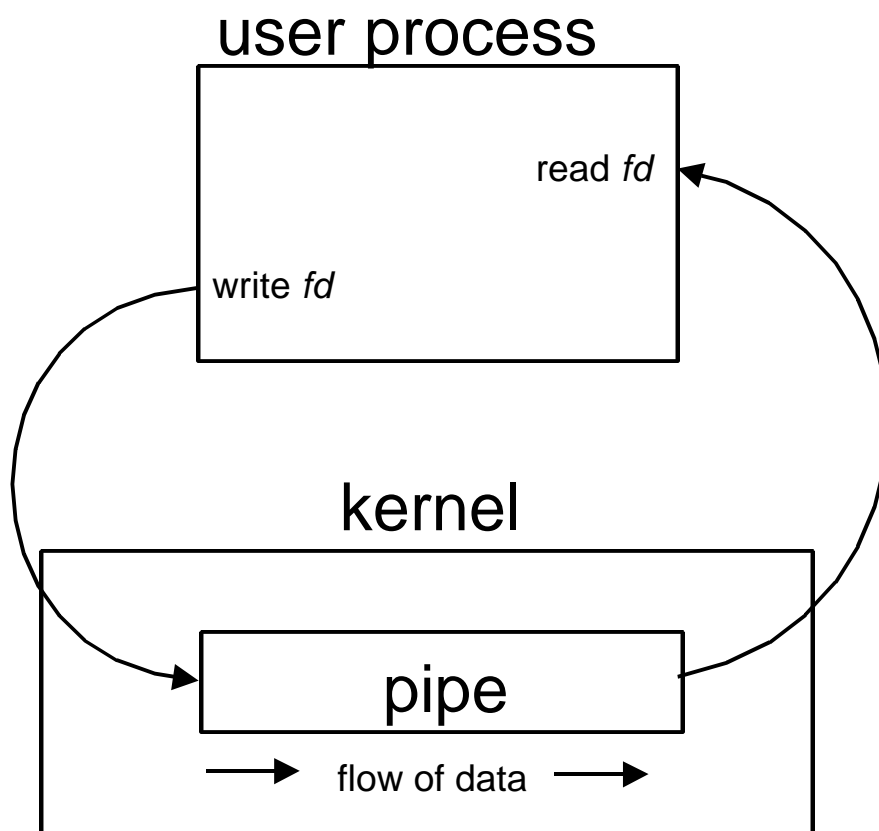
- Client writes filename to IPC channel.
- Server reads filename and
  - either returns contents via IPC
  - or returns opening error via IPC
- Client reads from IPC.



# Unix IPC - Pipes

## *pipe* System Call

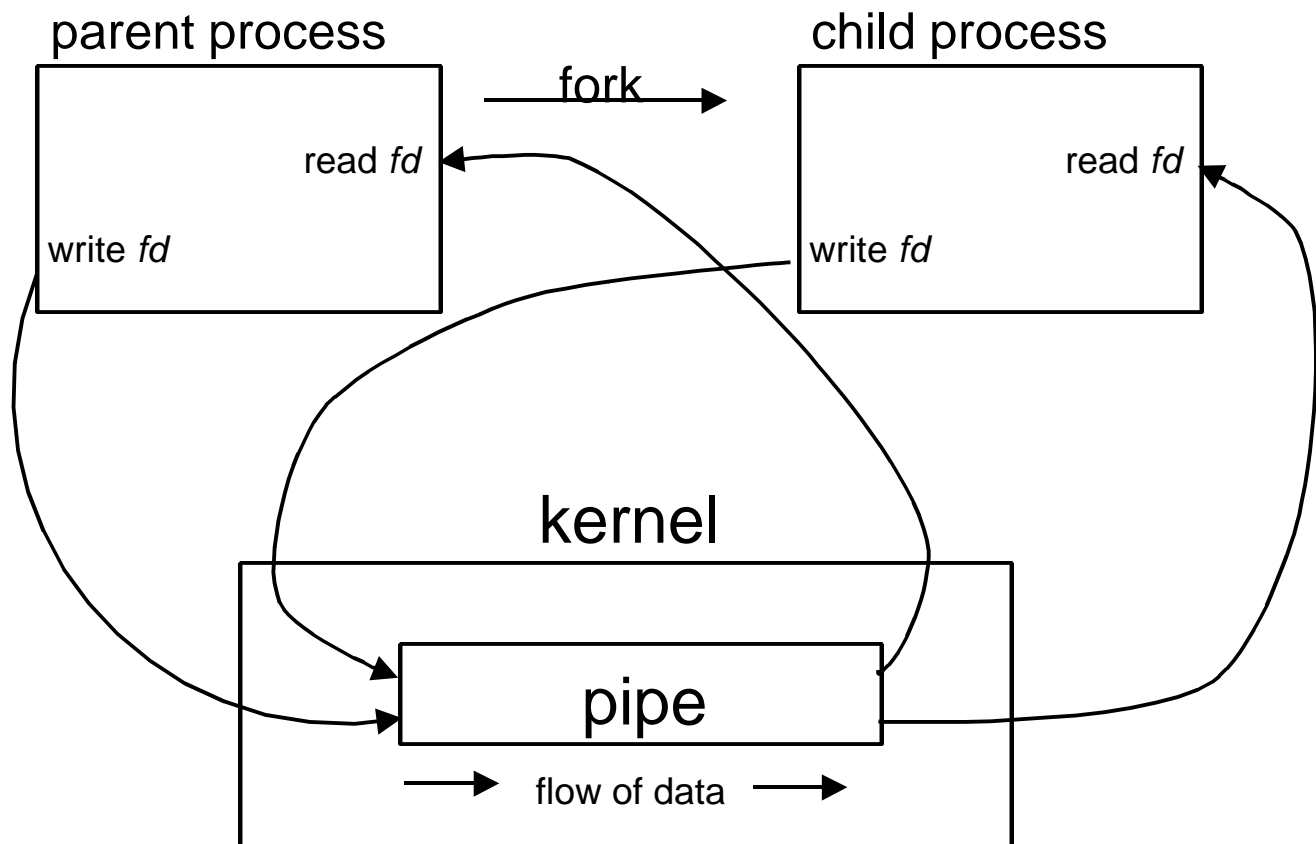
- One-way flow of data created by  
`int pipe( int *filedes);`
- Two file descriptors are returned  
*filedes*[0] open for reading  
*filedes*[1] open for writing



# Using a Pipe for IPC

Parent process forks after creating a pipe

- Open file descriptors copied in child
- Parent, child share both ends of pipe



Note: As an example, the parent process could be the producer and the child process could be the consumer.

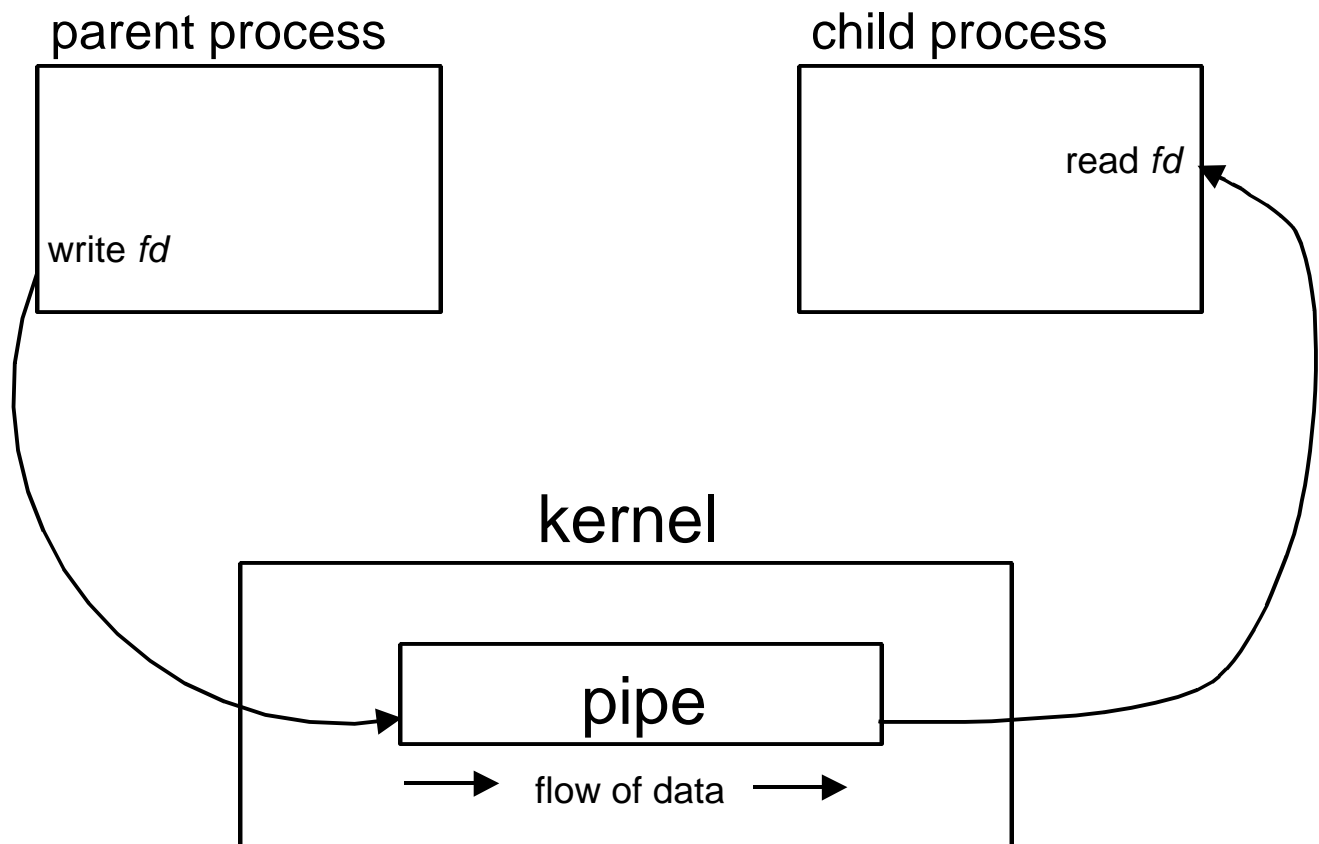
The producer should write to the pipe.

The consumer should read from the pipe.

Note: Pipes are only usable between processes that share an ancestor.

# Pipe Between Two Processes

After the *fork* that creates the shared pipe  
Writer closes the read end of pipe  
Reader closes the write end of pipe



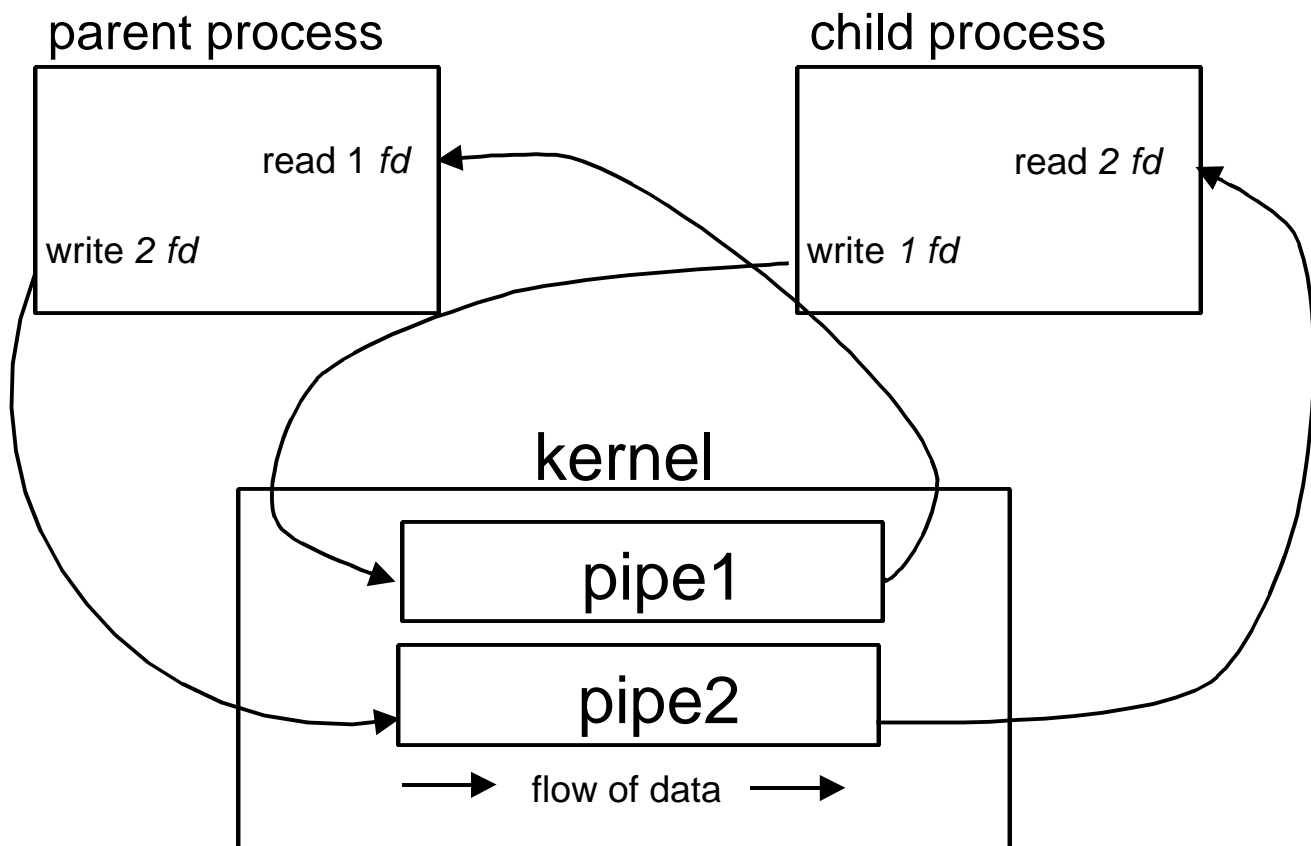
For 2-way communication  
Open 2 pipes before forking.  
Use 1 pipe as above.  
Reverse roles for 2nd pipe.

# Rules for Reading and Writing Pipes

- A *read* returns the lesser of the requested amount of data or contents of pipe.
- A *read* waits until there is data in the pipe if pipe is open for writing; else, it returns 0 for EOF.
- A *write* is atomic if data is less than the capacity of the pipe; else it is equivalent to many atomic *write* operations.
- A *write* waits until space is available if the pipe is full and open for reading; else the SIGPIPE signal is generated and 0 is returned. (If process does not catch this signal it is terminated.)
- The data is an uninterpreted stream of bytes.



# Client-Server Example



```
main()
{
    int child, pipe1[2], pipe2[2];
    if (pipe(pipe1) < 0 || pipe(pipe2) < 0) err_sys("can't create pipes");
    if ((child = fork()) < 0) err_sys("can't fork");
    else if (child > 0) { /* PARENT CODE */
        close(pipe1[1]); close(pipe2[0]);
        client(pipe1[0], pipe2[1]);
        while (wait((int *) 0) != child); /* WAIT FOR THIS CHILD */
        close(pipe1[0]); close(pipe2[1]);
        exit(0);
    }
    else { /* CHILD CODE */
        close(pipe1[0]); close(pipe2[1]);
        server(pipe2[0], pipe1[1]);
        close(pipe2[0]); close(pipe1[1]);
        exit(0);
    }
}
```

# Client Implementation

The client writes filename to the server.

- ★Read filename from *stdin*.
- ★Write to IPC channel.

Read data from server.

- ★Read IPC channel.
- ★Write to *stdout*.

```
#include <stdio.h>
#define MAXBUFF 1024
client(readfd, writefd)
int readfd; int writefd;
{
    char buff[MAXBUFF];
    int n;
    if (fgets(buff, MAXBUFF, stdin) == NULL)
        err_sys("client: filename read error");
    n = strlen(buff);
    if (buff[n-1] == '\n') n--;
    if (write(writefd, buff, n) != n)
        err_sys("client: filename write error");
    while ((n = read(readfd, buff, MAXBUFF)) > 0)
        if (write(STDOUT, buff, n) != n)
            err_sys("client: data write error");
        if (n < 0) err_sys("client: data read error");
}
```

# Server Implementation

- Read filename from IPC channel.
- If *open* fails write error on IPC.
- Else write data on IPC channel.

```
#include <stdio.h>
#define MAXBUFF 1024
server(readfd, writefd)
int readfd; int writefd;
{
    char buff[MAXBUFF], errmsg[256], *sys_err_str();
    int n, fd;
    extern int errno;
    if ((n = read(readfd, buff, MAXBUFF)) <= 0)
        err_sys("server: filename read error");
    buff[n] = '\0';
    if ( (fd = open(buff, 0) ) < 0) {
        sprintf(errmsg, ": can't open, %s\n", sys_err_str());
        strcat(buff, errmsg);
        n = strlen(buff);
        if (write(writefd, buff, n) != n)
            err_sys("server: errmsg write error");
    } else {
        while ((n = read(fd, buff, MAXBUFF)) > 0)
            if (write(writefd, buff, n) != n)
                err_sys("server: data write error");
            if (n < 0) err_sys("server: data read error");
    }
}
```

# Message Queue IPC (System V)

- System calls for message queue IPC.
  - msgget System call to create or open an already created message queue.
  - msgctl System call for control operations on message queue.
  - msgsnd System call for sending a message to message queue.
  - msgrcv System call to receive a message from message queue.
    - ➔ IPC key Used to identify the message queue. Created by calling `key_t ftok(char *pathname, char proj);` with `pathname` and `char` agreed upon by client and server.
  - <sys/msg.h> Include file for message queues.

# Client-Server Example using Message Queue IPC

## Message Data Structure (mesg.h)

```
typedef struct {
    int mesg_len;
    long mesg_type;
    char mesg_data[MAXMESGDATA];
} Mesg;
```

## Client Process

```
Mesg mesg;
client(ipcreadfd, ipcwritefd)
int ipcreadfd; int ipcwritefd;
{
    int n;
    if (fgets(mesg.mesg_data, MAXMSGDATA, stdin) == NULL)
        err_sys("Client: filename read error");
    n = strlen(mesg.mesg_data);
    if (mesg.mesg_data[n-1] == '\n') n--;
    mesg.mesg_len = n;
    mesg.mesg_type = 1L;
    mesg_send(ipcwritefd, &mesg);

    while ( (n = mesg_rcv(ipcreadfd, &mesg)) > 0)
        if (write(1, mesg.mesg_data, n) != n)
            err_sys("Client: Data write error");
        if (n < 0) err_sys("Client: Data read error");
}
```

# Client-Server Example (continued)

## Server Process

```
Mesg mesg;
server(ipcreadfd, ipcwritefd)
int ipcreadfd; int ipcwritefd;
{
    int n, fileid;
    char errmesg[256], *sys_err_str();

    mesg.mesg_type = 1L;
    if ( (n = mesg_rcv(ipcreadfd, &mesg)) <= 0)
        err_sys("Server: Filename read error");
    mesg.mesg_data[n] = '\0';
    if ( (fileid = open(mesg.mesg_data, 0)) < 0) {
        sprintf(errmesg, ": can't open, %s\n", sys_err_str());
        strcat(mesg.mesg_data, errmesg);
        mesg.mesg_len = strlen(mesg.mesg_data);
        mesg_send(ipcwritefd, &mesg);
    }
    else {
        while ( (n = read(fileid, mesg.mesg_data, MAXMESGDATA)) > 0)
        {mesg.mesg_len = n;
         mesg_send(ipcwritefd, &mesg);
        }
        close(fileid);
        if (n < 0) err_sys("Server: Data read error");
    }
    /* Empty message to denote end */
    mesg.mesg_len = 0;
    mesg_send(ipcwritefd, &mesg);
}
```

# Client-Server main Functions

```
#include "msgq.h"
```

```
main()
```

```
{
```

```
    int readfd, writefd;
```

```
    if ((readfd = msgget(MKEY1, PERMS | IPC_CREAT) ) < 0)
        err_sys("Server: Can't get message queue 1");
```

```
    if ((writefd = msgget(MKEY2, PERMS | IPC_CREAT) ) < 0)
        err_sys("Server: Can't get message queue 2");
```

```
    server(readfd, writefd);
```

```
    exit(0);
```

```
}
```

```
#include <sys/types.h>
#include <sys/ipc.h>
#include <sys/msg.h>
```

```
#include <sys/errno.h>
extern int errno;
```

```
#define MKEY1 1234L
#define MKEY2 2345L
```

```
#define PERMS 0666
```

## msgq.h

```
#include "msgq.h"
```

```
main()
```

```
{
```

```
    int readfd, writefd;
```

```
    if ((writefd = msgget(MKEY1, PERMS | IPC_CREAT) ) < 0)
        err_sys("Client: Can't get message queue 1");
```

```
    if ((readfd = msgget(MKEY2, PERMS | IPC_CREAT) ) < 0)
        err_sys("Client: Can't get message queue 2");
```

```
    client(readfd, writefd);
```

```
    if (msgctl(readid, IPC_RMID, (struct msqid_ds *) 0) ) < 0)
        err_sys("Client: Can't RMID message queue 2");
```

```
    if (msgctl(writeid, IPC_RMID, (struct msqid_ds *) 0) ) < 0)
        err_sys("Client: Can't RMID message queue 1");
```

```
    exit(0);
```

```
}
```

# Data Abstraction

## mesg\_send & mesg\_rcv

```
#include "mesg.h"
```

```
mesg_send(msgQid, msgptr)
```

```
int msgQid;
```

```
Mesg *msgptr;
```

```
{
```

```
    /* Use msgsnd system call to send message */
```

```
    if (msgsnd(msgQid, (char *) &(msgptr->mesg_type),  
        msgptr->mesg_len, 0) != 0)  
        err_sys("msgsnd error");
```

```
}
```

```
#include "mesg.h"
```

```
int mesg_rcv(msgQid, msgptr)
```

```
int msgQid;
```

```
Mesg *msgptr;
```

```
{
```

```
    int n;
```

```
    /* Read the first message of specified message type */
```

```
    n = msgrcv(msgQid, (char *) &(msgptr->mesg_type),  
        MAXMSGDATA, msgptr->mesg_type, 0);
```

```
    if ( (msgptr->mesg_len = n) < 0)
```

```
        err_dump("msgrcv error");
```

```
    return (n);        /* n = 0 signifies EOF */
```

```
}
```